

Catalytic Compilation

A modern take on safe Spark









- Distributed relational engine since 2014
- Scala's Killer App
- Large ecosystem : data sources, languages

```
SELECT s.name, r.timestamp, r.temperature
FROM weather_stations s
JOIN weather_readings r ON s.id = r.station_id
WHERE r.temperature > 30.0
ORDER BY r.temperature DESC
```

SQL String

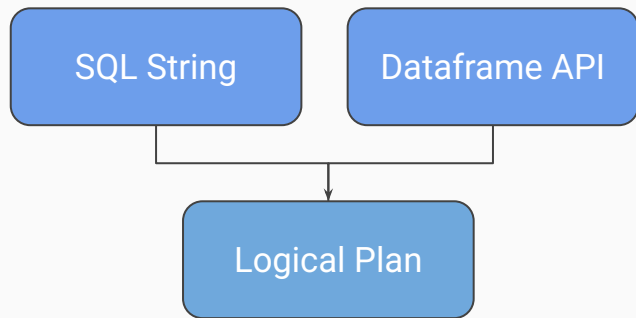
Dataframe API

```
val df = spark.sql:  
    """  
    SELECT s.name, r.timestamp, r.temperature  
    FROM weather_stations s  
    JOIN weather_readings r ON s.id = r.station_id  
    WHERE r.temperature > 30.0  
    ORDER BY r.temperature DESC  
    """
```

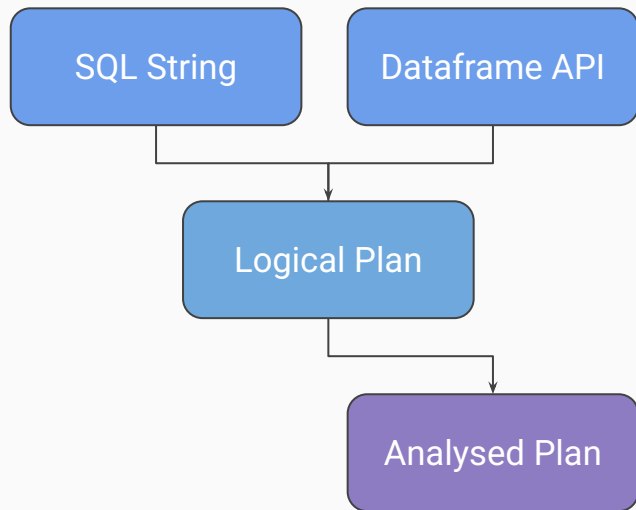
SQL String

Dataframe API

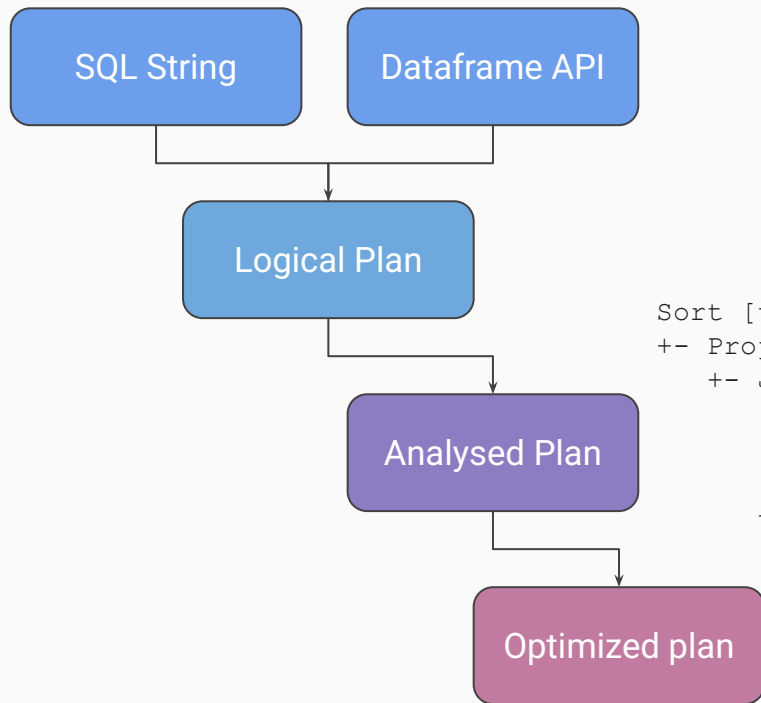
```
stations
    .as("s")
    .join(readings.as("r"), stations.col("id") === readings.col("station_id"), "inner")
    .where(col("r.temperature") > 30.0)
    .select(
        col("s.name").as("name"),
        col("r.timestamp").as("timestamp"),
        col("r.temperature").as("temperature")
    )
    .orderBy(col("r.temperature").desc)
```



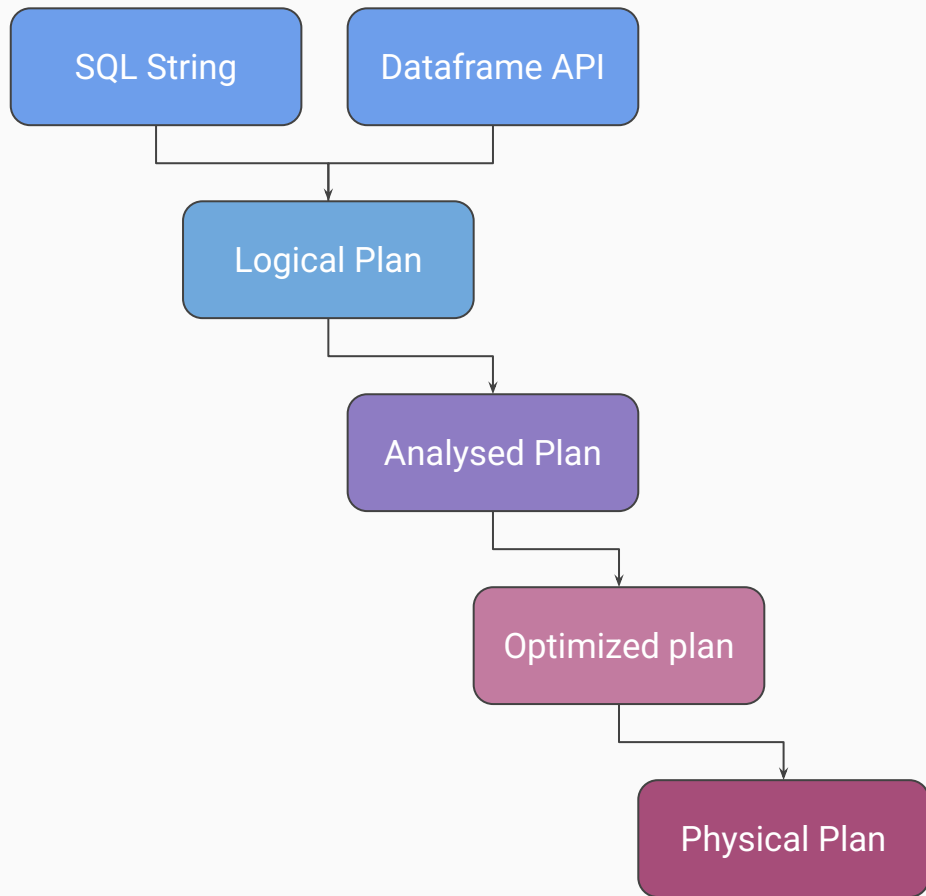
```
'Sort ['r.temperature DESC NULLS LAST], true
+- 'Project ['s.name, 'r.timestamp, 'r.temperature]
  +- 'Filter ('r.temperature > 30.0)
    +- 'Join Inner, ('s.id = 'r.station_id)
      :- 'SubqueryAlias s
      : +- 'UnresolvedRelation [weather_stations], [], false
    +- 'SubqueryAlias r
      +- 'UnresolvedRelation [weather_readings], [], false
```

```
Sort [temperature#6 DESC NULLS LAST], true
+- Project [name#1, timestamp#5, temperature#6]
  +- Filter (temperature#6 > cast(30.0 as double))
    +- Join Inner, (id#0 = station_id#4)
      :- SubqueryAlias s
      :   +- SubqueryAlias compiletime.default.weather_stations
      :     +- RelationV2[id#0, name#1, latitude#2, longitude#3] wea
    +- SubqueryAlias r
      +- SubqueryAlias compiletime.default.weather_readings
        +- RelationV2[station_id#4, timestamp#5, ...] weather_re
```



```
Sort [temperature#27 DESC NULLS LAST], true
+- Project [name#11, timestamp#28, temperature#27]
  +- Join Inner, (id#8L = station_id#26L)
    :- Project [id#8L, name#11]
    :   +- Filter isnonnull(id#8L)
    :     +- Relation [id#8L,latitude#9,longitude#10,name#11] json
  +- Project [station_id#26L, temperature#27, timestamp#28]
    +- Filter ((isnonnull(temperature#27) AND (temperature#27 > 30
      AND isnonnull(station_id#26L))
    +- Relation [humidity#24,pressure#25,station_id#26L,tempera
      timestamp#28] json
```



```

AdaptiveSparkPlan isFinalPlan=false
+- Sort [temperature#27 DESC NULLS LAST], true, 0
  +- Exchange rangepartitioning(temperature#27 DESC NULLS LAST, 200), ...
    +- Project [name#11, timestamp#28, temperature#27]
      +- BroadcastHashJoin [id#8L], [station_id#26L], Inner, BuildLeft, false
        :- BroadcastExchange [params...]
        :   +- Filter isnotnull(id#8L)
        :     +- FileScan json [id#8L,name#11]
        :       Batched: false,
        :       DataFilters: [isnotnull(id#8L)],
        :       Format: JSON, Location: InMemoryFileIndex(1 paths)[file],
        :       PartitionFilters: [],
        :       PushedFilters: [IsNotNull(id)],
        :       ReadSchema: struct<id:bigint,name:string>
      +- Filter (filters...)
        +- FileScan json [station_id#26L,temperature#27,timestamp#28]
          Batched: false,
          DataFilters: [filters],
          Format: JSON,
          Location: InMemoryFileIndex(1 paths)[file],
          PartitionFilters: [],
          PushedFilters: [filters],
          ReadSchema: struct<station_id:bigint,temperature:double,timestamp:string>

```



Catalyst is just a compiler



Catalyst is just a compiler *in a runtime*



Catalyst is just a compiler *in a runtime **without tooling***

Strongly typed pipelines ?

```
val df = spark.sql:  
    """  
    SELECT s.name, r.timestamp r.temperature  
    FROM weather_stations s  
    JOIN weather_readings r ON s.id = r.station_id  
    WHERE r.temperature > 30.0  
    ORDER BY r.temperature DESC  
    """
```


Strongly typed pipelines ?

```
[PARSE_SYNTAX_ERROR] Syntax error at or near '.'. (line 2, pos 30)
```

```
== SQL ==
```

```
SELECT s.name, r.timestamp r.temperature
-----^^^
FROM weather_stations s
JOIN weather_readings r ON s.id = r.station_id
WHERE r.temperature > 30.0
ORDER BY r.temperature DESC
```

Strongly typed pipelines ?

```
stations
  .as("s")
  .join(readings.as("r"), stations.col("id") === readings.col("station_id"), "inner")
  .where(col("r.temperature") > 30.0)
  .select(
    col("s.name").as("name"),
    col("r.timestamp").as("timestamp"),
    col("r.temprature").as("temperature")
  )
  .orderBy(col("r.temperature").desc)
```

Strongly typed pipelines ?

```
[UNRESOLVED_COLUMN.WITH_SUGGESTION] A column or function parameter with name `r`.`temprature` cannot be resolved. Did you mean one of the following? [`r`.`temperature`, `r`.`pressure`, `r`.`humidity`, `s`.`latitude`, `s`.`longitude`].;
```

```
'Project [name#11 AS name#61, timestamp#28 AS timestamp#62, 'r.temperature AS temperature#63]
```

```
  +- Filter (temperature#27 > 30.0)
```

```
    +- Join Inner, (id#8L = station_id#26L)
```

```
      :- SubqueryAlias s
```

```
      :   +- SubqueryAlias weather_stations
```

```
      :     +- View (`weather_stations`, [id#8L,latitude#9,longitude#10,name#11])
```

```
      :       +- Relation [id#8L,latitude#9,longitude#10,name#11] json
```

```
    +- SubqueryAlias r
```

```
      +- SubqueryAlias weather_readings
```

```
      +- View (`weather_readings`, [...])
```

```
      +- Relation [humidity#24,pressure#25,station_id#26L,temperature#27,timestamp#28] json
```

Stringly typed pipelines !



Why did pyspark succeed ?

```
from pyspark.sql.functions import col

stations = spark.table("weather_stations")
readings = spark.table("weather_readings")

df = (
    stations.alias("s")
    .join(
        readings.alias("r"),
        col("s.id") == col("r.station_id"),
        "inner"
    )
    .where(col("r.temperature") > 30.0)
    .select(
        col("s.name").alias("name"),
        col("r.timestamp").alias("timestamp"),
        col("r.temperature").alias("temperature")
    )
    .orderBy(col("r.temperature").desc())
)
```

```
import org.apache.spark.sql.functions.*

val stations = spark.table("weather_stations")
val readings = spark.table("weather_readings")

val df =
    stations.as("s")
    .join(
        readings.as("r"),
        stations.col("id") === readings.col("station_id"),
        "inner"
    )
    .where(col("r.temperature") > 30.0)
    .select(
        col("s.name").as("name"),
        col("r.timestamp").as("timestamp"),
        col("r.temperature").as("temperature")
    )
    .orderBy(col("r.temperature").desc())
```

Why did pyspark succeed ?



The issue with Spark SQL

- Scala add **little to no value** in term of correctness
- Scala compiler do not understand table and rows
- All checks are deferred until runtime (or tests)

And then ?

What if there was a better way ?

typelevel/**frameless**

Expressive types for Spark.



 54

Contributors

 37

Issues

 889

Stars

 137

Forks



Frameless

```
case class WeatherStation(id: Int, name: String, latitude: Double, longitude: Double)
case class WeatherReading(station_id: Int, timestamp: java.sql.Timestamp, temperature: Double, humidity: Double, pressure: Double)

implicit val stationEnc: Encoder[WeatherStation] = Encoders.product
implicit val readingEnc: Encoder[WeatherReading] = Encoders.product

spark.read.json("./stations.json").createTempView("weather_stations")
spark.read.json("./readings.json").createTempView("weather_readings")

val stations = TypedDataset.create(spark.table("weather_stations").as[WeatherStation])
val readings = TypedDataset.create(spark.table("weather_readings").as[WeatherReading])

@nowarn
val joined = stations.joinInner(readings)(stations('id) === readings('station_id))

@nowarn
val selected = joined
  .filter(joined('_2).field('temperature) > 30.0)
  .select(
    joined('_1).field('name),
    joined('_2).field('timestamp),
    joined('_2).field('temperature)
  )

@nowarn
val df = selected.orderBy(selected('_3).desc)
```

Frameless

```
case class WeatherStation(id: Int, name: String, latitude: Double, longitude: Double)
case class WeatherReading(station_id: Int, timestamp: java.sql.Timestamp, temperature: Double, humidity: Double, pressure: Double)

implicit val stationEnc: Encoder[WeatherStation] = Encoders.product
implicit val readingEnc: Encoder[WeatherReading] = Encoders.product

spark.read.json("./stations.json").createTempView("weather_stations")
spark.read.json("./readings.json").createTempView("weather_readings")

val stations = TypedDataset.create(spark.table("weather_stations").as[WeatherStation])
val readings = TypedDataset.create(spark.table("weather_readings").as[WeatherReading])

@nowarn
val joined = stations.joinInner(readings)(stations('id) === readings('station_id))

@nowarn
val selected = joined
  .filter(joined('_2).field('temperature) > 30.0)
  .select(
    joined('_1).field('name),
    joined('_2).field('timestamp),
    joined('_2).field('temperature)
  )

@nowarn
val df = selected.orderBy(selected('_3).desc)
```



Uses Shapeless
Witnesses

Frameless

```
case class WeatherStation(id: Int, name: String, latitude: Double, longitude: Double)
case class WeatherReading(station_id: Int, timestamp: java.sql.Timestamp, temperature: Double, humidity: Double, pressure: Double)

implicit val stationEnc: Encoder[WeatherStation] = Encoders.product
implicit val readingEnc: Encoder[WeatherReading] = Encoders.product

spark.read.json("./stations.json").createTempView("weather_stations")
spark.read.json("./readings.json").createTempView("weather_readings")

val stations = TypedDataset.create(spark.table("weather_stations").as[WeatherStation])
val readings = TypedDataset.create(spark.table("weather_readings").as[WeatherReading])

@nowarn
val joined = stations.joinInner(readings)(stations('id) === readings('station_id))

@nowarn
val selected = joined
  .filter(joined('_2).field('temperature) > 30.0)
  .select(
    joined('_1).field('name),
    joined('_2).field('timestamp),
    joined('_2).field('temperature)
  )

@nowarn
val df = selected.orderBy(selected('_3).desc)
```



Uses Shapeless
Witnesses



Quoted Symbols
Deprecation in 2.13

Frameless

```
case class WeatherStation(id: Int, name: String, latitude: Double, longitude: Double)
case class WeatherReading(station_id: Int, timestamp: java.sql.Timestamp, temperature: Double, humidity: Double, pressure: Double)

implicit val stationEnc: Encoder[WeatherStation] = Encoders.product
implicit val readingEnc: Encoder[WeatherReading] = Encoders.product

spark.read.json("./stations.json").createTempView("weather_stations")
spark.read.json("./readings.json").createTempView("weather_readings")

val stations = TypedDataset.create(spark.table("weather_stations").as[WeatherStation])
val readings = TypedDataset.create(spark.table("weather_readings").as[WeatherReading])

@nowarn
val joined = stations.joinInner(readings)(stations('id) === readings('station_id))

@nowarn
val selected = joined
  .filter(joined('_2).field('temperature) > 30.0)
  .select(
    joined('_1).field('name),
    joined('_2).field('timestamp),
    joined('_2).field('temperature)
  )

@nowarn
val df = selected.orderBy(selected(_3).desc)
```



Uses Shapeless
Witnesses



Quoted Symbols
Deprecation in 2.13

No column Symbol with shapeless.tag.Tagged[String("temperature")] of type V in Playground.App.WeatherReading

And then ?

What if there was a better way ? (season 2)

VirtusLab/iskra

Typesafe wrapper for Apache Spark DataFrame API



 3

Contributors

 7

Issues

 141

Stars

 9

Forks



```
// no iska encoder defined for java.sql.Timestamp

case class WeatherStation(id: Int, name: String, latitude: Double, longitude: Double)
case class WeatherReading(station_id: Int, timestamp: Long, temperature: Double, humidity: Double, pressure: Double)

spark.read.json("./stations.json").createTempView("weather_stations")
spark.read.json("./readings.json").createTempView("weather_readings")

val stations = spark.table("weather_stations").typed[WeatherStation]
val readings = spark.table("weather_readings").typed[WeatherReading]

val df = stations
  .innerJoin(readings)
  .on($.stations.id === $.readings.station_id)
  .where($.readings.temperature > lit(30.0))
  .select(
    $.stations.name,
    $.readings.timestamp,
    $.readings.temperature
  )
// .orderBy($.readings.temperature.desc)
```



```
// no iska encoder defined for java.sql.Timestamp

case class WeatherStation(id: Int, name: String, latitude: Double, longitude: Double)
case class WeatherReading(station_id: Int, timestamp: Long, temperature: Double, humidity: Double, pressure: Double)

spark.read.json("./stations.json").createTempView("weather_stations")
spark.read.json("./readings.json").createTempView("weather_readings")

val stations = spark.table("weather_stations").typed[WeatherStation]
val readings = spark.table("weather_readings").typed[WeatherReading]

val df = stations
  .innerJoin(readings)
  .on($.stations.id === $.readings.station_id)
  .where($.readings.temperature > lit(30.0))
  .select(
    $.stations.name,
    $.readings.timestamp,
    $.readings.temperature
  )
// .orderBy($.readings.temperature.desc)
```



Some rough edges

```
// no iska encoder defined for java.sql.Timestamp

case class WeatherStation(id: Int, name: String, latitude: Double, longitude: Double)
case class WeatherReading(station_id: Int, timestamp: Long, temperature: Double, humidity: Double, pressure: Double)

spark.read.json("./stations.json").createTempView("weather_stations")
spark.read.json("./readings.json").createTempView("weather_readings")

val stations = spark.table("weather_stations").typed[WeatherStation]
val readings = spark.table("weather_readings").typed[WeatherReading]

val df = stations
  .innerJoin(readings)
  .on($.stations.id === $.readings.station_id)
  .where($.readings.temperature > lit(30.0))
  .select(
    $.stations.name,
    $.readings.timestamp,
    $.readings.temperature
  )
// .orderBy($.readings.temperature.desc)
```



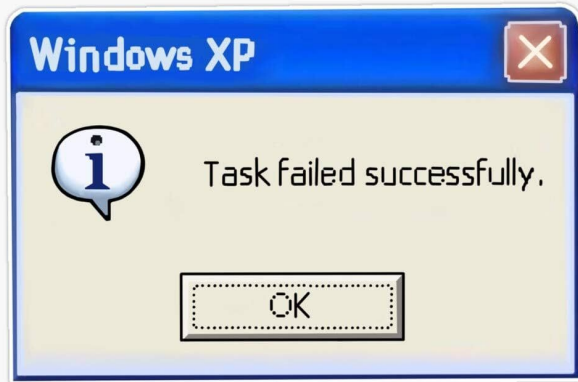
Some rough edges



Use its own Codecs

And then ?

Brute forcing the problem



- High migration costs
- Everything needs to be modeled
- The ORM syndrome : another tool to learn

And then ?

What if there was a better way ? (season 3)

And then ?

Let's put spark into the compiler !



And then ?

Live demo

Spark Compiletime

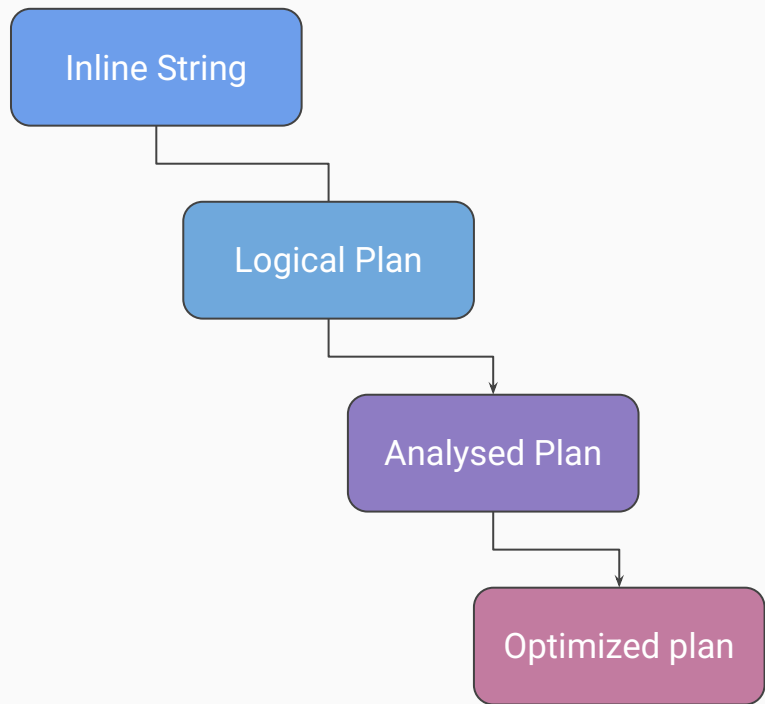
```
[Analyzer] : [UNRESOLVED_COLUMN_FOR_JOIN] USING column `station_id` cannot be resolved on the left side of the join. The left-side columns: [`id`, `latitude`, `longitude`, `name`]. sbt
```

[View Problem \(Alt+F8\)](#) [Quick Fix... \(Ctrl+.\)](#) [Fix using Copilot \(Ctrl+I\)](#)

```
val df = spark.sql(weatherDb):  
    """  
    SELECT s.name, r.timestamp, r.temperature  
    FROM weather stations s  
    JOIN weather readings r using (station_id)  
    WHERE r.temperature > 30.0  
    ORDER BY r.temperature DESC  
    """
```

Using macros, Catalyst itself is parsing and optimizing the query

How does it work ?



- Type level catalog using mirrors
- Table mirrors are derived from their creation query
- Macros runs lightweight catalog (no SparkSession)

Why tables mirrors

Table mirrors : from inlines ...

```
trait TableMirror:  
  type DB <: String  
  type Name <: String  
  type Schema <: String  
  type Query <: String  
  
  inline def db: String          = TableMirror.db[this.type]  
  inline def name: String        = TableMirror.name[this.type]  
  inline def schema: StructType = TableMirror.schema[this.type]  
  inline def schemaString: String = TableMirror.schemaString[this.type]  
  inline def query: String       = TableMirror.query[this.type]
```

Table mirrors : to macros ...

```
object TableMirror:

  inline def query[T]: String =
    ${ macros.queryImpl[T] }

  // other implementations
```

Table mirrors : to type matching

```
def queryImpl[T: Type](using Quotes): Expr[String] =  
  import quotes.reflect.*  
  Type.of[T] match  
    case '[TableMirror { type Query = query & String }] =>  
      Expr(utils.stringFromType[query])
```

Table mirrors : to AST handling

```
def stringFromType[T: Type](using Quotes): String =  
  import quotes.reflect.*  
  TypeRepr.of[T] match  
    case ConstantType(StringConstant(label)) => label  
    case _ =>  
      report.errorAndAbort(s"expected a constant string, got ${TypeRepr.of[T]}")
```

Table mirrors : to AST handling

```
def typeFromString(name: String) (using Quotes): Type[?] =  
  import quotes.reflect.*  
  ConstantType(StringConstant(name)).asType
```

Table mirrors : materialization

```
val dbType      = utils.typeFromString(tableName.namespace().mkString("."))
val nameType    = utils.typeFromString(tableName.name())
val schemaType  = utils.typeFromString(tableSchema.toDDL)
val queryType   = utils.typeFromString(sql)
```

```
(dbType, nameType, schemaType, queryType) match
```

```
case ('[db]', '[name]', '[schema]', '[query]) =>
```

```
'{
```

```
  new TableMirror {
```

```
    type DB      = db      & String
```

```
    type Name    = name    & String
```

```
    type Schema  = schema  & String
```

```
    type Query   = query   & String
```

```
  }
```

```
}
```

```
case unreachable =>
```

```
  report.errorAndAbort(s"Unexpected types: $unreachable")
```

Toward the future

- Table mirrors are only extractable from SQL statement
 - Other sources could be used : avrodl, DB schemas, Hive catalogs
- No concatenation / interpolators
- API is rough while interacting with datasets / dataframes

Conclusion

- Embrace SQL as your relational data manipulation DSL
- Scala macros are super fun to use !
- Scala shines at making compiler extensions
 - ideas: Apache calcite, datalog, or more (CoQ?)

Thank you !