



Un Éléphant de Première Classe

A bit of History

1970

A Relational Model of Data for Large Shared Data Banks

E. F. Codd

IBM Research Laboratory, San Jose, California

1970 - 1979 System R



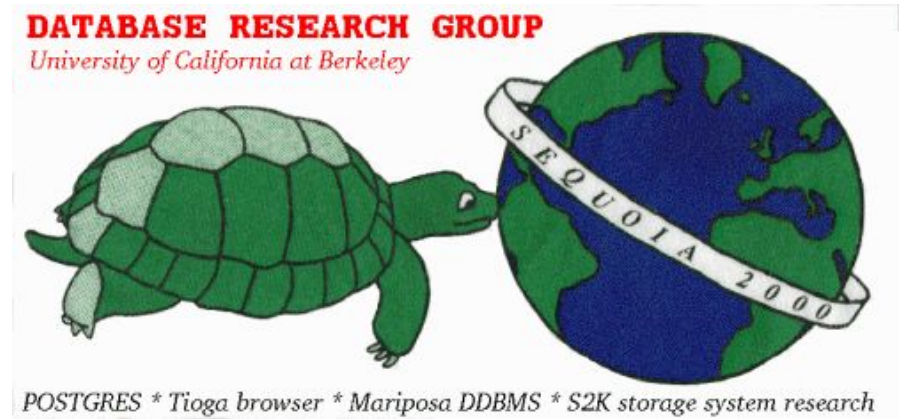
1973 - 1985

Ingres & Quel



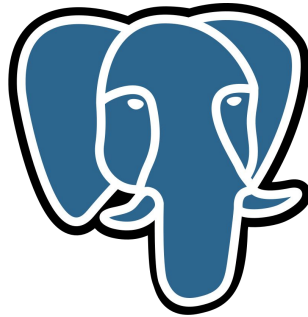
1989

Postgres & PostQuel



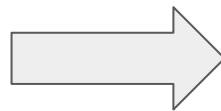
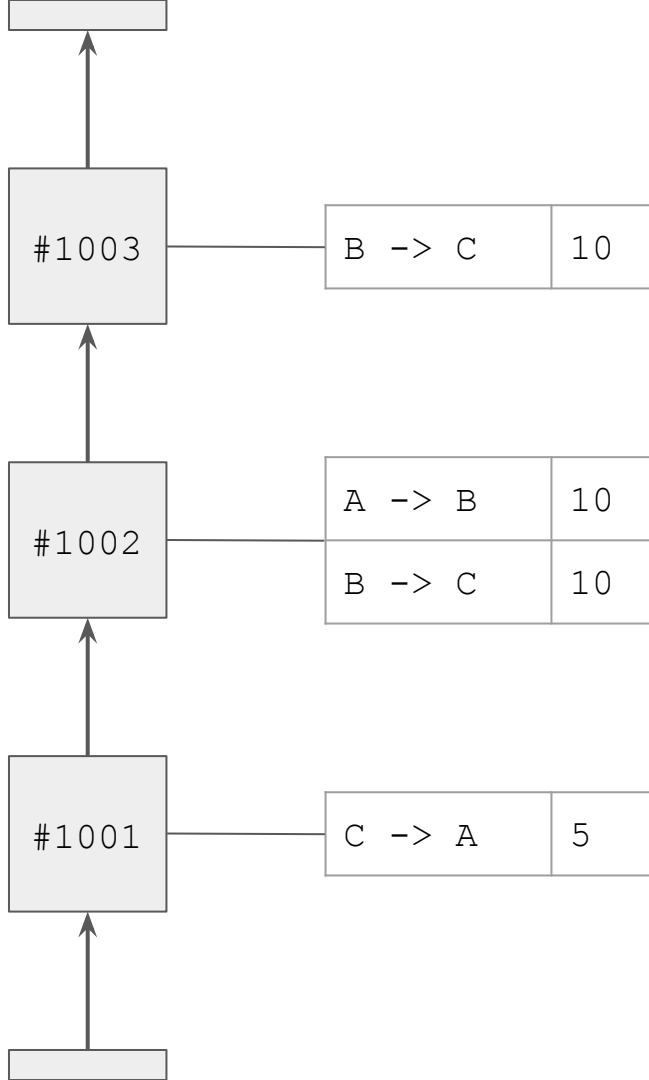
1996

PostgreSQL



Why this talk ?

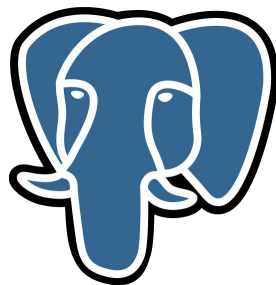
[LEDGER]



A	1001	+5
A	1002	-10
B	1001	+10
B	1002	+10
C	1001	-10
C	1002	-10



+





✗ Optimized for write-based workflow



- ✗ Optimized for write-based workflow
- ✗ Heavy in term of maintenance



- ✗ Optimized for write-based workflow
- ✗ Heavy in term of maintenance
- ✗ Joins must be performed on the application side

Applicative joins ?

```
class OrderService:

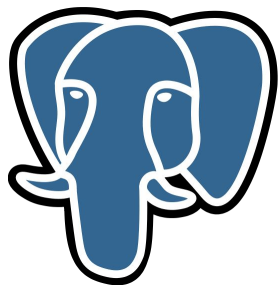
    val orderRepository: OrderRepository = ???
    val customerRepository: CustomerRepository = ???

    def findLatestOrdersWithCustomer(): Seq[OrderWithCustomer] =
        val latests = orderRepository.findLatestOrders()

        val result = ListBuffer.empty[OrderWithCustomer]

        latests.foreach: order =>
            val customer = customerRepository.findById(order.customerID)
            result.addOne(OrderWithCustomer(order, customer))

        result.toList
```

- ✓ Great for our read-heavy workflow
- ✓ Lighter maintenance and cheaper (50%)
- ✓ Joins can just be joins !

What is first class Postgres ?

- Database is at the center of the system
- The schema is the source of truth
- Treat SQL as a *real programming language*
- Commit to Postgres features

Setup your development loop

Tools for composition

Step up your indexing game

Data types shenanigans

Setup your development loop

Tools for composition

Step up your indexing game

Data types shenanigans

Hardware Setup



- Modern hardware is worth it
- Scaling vertically is a viable strategy
- On-prem kube (rancher) + bare metal PG

Dual Intel Xeon Gold 5515+ - 16c/32t - 3.2GHz/3.6GHz

256GB DDR5 ECC 4800MHz

2x SSD NVMe 960GB Datacenter Class Soft RAID

6x 3.84TB SSD NVMe Soft RAID

- Modern hardware is worth it
- Scaling vertically is a viable strategy
- On-prem kube (rancher) + bare metal PG

Dual Intel Xeon Gold 5515+ - 16c/32t - 3.2GHz/3.6GHz

256GB DDR5 ECC 4800MHz

2x SSD NVMe 960GB Datacenter Class Soft RAID

6× 3.84TB SSD NVMe Soft RAID

~ 400 €

Software Setup

Use a migration tool (we use flyway)

Prefer non-idempotent migrations

Think about scheduling baselines

Abuse comments on everything !



Abuse comments on everything !

```
-- Transaction => block join table, a.k.a. confirmation (inclusion of a transaction in a block)
create table v2.chain_tx_block
(
    block_hash text not null,
    tx_id      text not null,

    unique (tx_id),
    primary key (block_hash, tx_id),
    foreign key (tx_id) references v2.chain_tx (tx_id),
    foreign key (block_hash) references v2.chain_block (block_hash) on delete cascade
);

comment on table v2.chain_tx_block is 'represents the inclusion of a transaction in a block (a.k.a., a
confirmation) - note that it can be a minority fork block';

comment on column v2.chain_tx_block.tx_id is 'transaction identifier (immutable, globally unique identifier,
typically transaction hash)';

comment on column v2.chain_tx_block.block_hash is 'block identifier (immutable, globally unique)';
```

No ORM : prefer data access libraries

skunk.exception.PostgresErrorException:

Postgres ERROR 42804 raised in transformAssignedExpr (parse_target.c:595)

Problem: Column "success" is of type boolean but expression is of type bigint.

Hint: You will need to rewrite or cast the expression.

The statement under consideration was defined

at /Users/raphael.lemaitre/Developer/projects/ledger/atlas/modules/polkadot/src/main/scala/polkadot/queries/ExtrinsicTable.scala:18

```
INSERT INTO extrinsics
  SELECT height                AS height,
         index                 AS index,
         data → 'hash'         AS hash,
         data → 'method' → 'pallet' AS pallet,
         data → 'method' → 'method' AS method,
         (data → 'nonce') :: bigint AS nonce,
         data → 'signature' → 'signer' → 'id' AS signer,
         data → 'args'                AS args,
         (data → 'tip') :: bigint      AS tip,
         (data → 'info' → 'weight') :: bigint AS weight,
         data → 'info' → 'class'       AS class,
         (data → 'info' → 'partialFee') :: bigint AS partial_fee,
         data → 'era'                  AS era,
         (data → 'success') :: bigint   AS success,
         └─ Column "success" is of type boolean but expression is of type bigint.
         (data → 'paysFee') :: boolean  AS pays_fee
  FROM raw_extrinsics
  WHERE height ≥ $1 AND height ≤ $2
```

If this is an error you wish to trap and handle in your application, you can do so with a `SqlState` extractor. For example:

```
doSomething.recoverWith { case SqlState.DatatypeMismatch(ex) ⇒ ... }
```

Use a proper SQL editor (DataGrip, DBeaver)

`psql` can be a very powerful tool

psql tips :

- `\e` Open the EDITOR with the latest query you typed. Closing run the buffer.
- `\e myfile.sql` Open the EDITOR with the file in it.
- `\x` For extended display
- `\watch 1` Repeat the query every 1 second

Testing SQL

Testing with a real DB

- Testcontainers + fixtures to add migration
- Cache DB instances !
- Never use sqlite / h2 / derby anymore



Crafting SQL Tests

- Test on the boundaries.
 - Typed language -> Input / Output models then translate into domain
 - Dynamic languages -> check aggressively
 - Defense in depth strategies
- Test functions and procedures
- Make tooling around sql files (syntax highlighting)

Testing on real data

Use transactions to benchmark your migrations before running them with migrations

```
begin isolation level read committed; -- Queries sees data before it starts
begin isolation level repeatable read; -- Queries sees data before the tx started
begin isolation level serializable;    -- repeatable read + error if write conflict

-- Migrate, migrate, migrate

savepoint before_risky_thing;

-- Migrate, migrate, migrate

-- Oh no ! I deleted the prod table :/

rollback to before_risky_thing; -- Rollback to savepoint, restoring the state

-- Migrate, migrate, migrate

commit; -- or abort if you want your migrations to pick it up
```

Optimizing

EXPLAIN (ANALYZE, COSTS, VERBOSE, BUFFERS)

```
explain (analyse, verbose, costs, buffers)
select header_id, tx_index, address
from log
      join address on contract_id = address_id
where (header_id, tx_index) = (2278166300, 290);
```

QUERY PLAN

```
Nested Loop  (cost=[...] rows=10 width=33) (actual time=0[...] rows=5 loops=1)
  Output: log.header_id, log.tx_index, address.address
  Inner Unique: true
  Buffers: shared hit=30
  -> Index Scan using log_pkey on v2.log  (cost=[...]) (actual time=[...] rows=5 loops=1)
      Output: [columns]
      Index Cond: ((log.header_id = '2278166300'::bigint) AND (log.tx_index = 290))
      Buffers: shared hit=5
  -> Index Scan using address_pkey on v2.address  (cost=[...]) (actual time=[...] rows=1 loops=5)
      Output: address.address_id, address.address
      Index Cond: (address.address_id = log.contract_id)
      Buffers: shared hit=25
Query Identifier: -8177880399755634493
Planning:
  Buffers: shared hit=10
Planning Time: 0.467 ms
Execution Time: 0.094 ms
(17 rows)
```

Always analyse after large inserts

Scans are the name of the game

- Seq Scan Traverse in order : maybe an index is missing ?
- Index Scan Scans a table heap using an index
- Index Only Scan Only touches the index pages
- Bitmap Index Scan Build a bitmap from an index
- Bitmap Heap Scan Build a bitmap from the heap
- TID Scan Directly access physical data

Sort is a very costly operation

- `union, distinct` introduce a sort
- Cover the order with the indices (and preserve order through the query)
- Sorting large set can spill on disk : `set work_mem` can help

Shared Hit data is coming from cached pages

Read data is coming from disk

Rows Removed by Filter means your indices are not efficient enough

explain.dalibo.com

Setup your development loop

Tools for composition

Step up your indexing game

Data types shenanigans

```
select first_name
       , last_name
       , performance_rating
       , project_count
from (select employee_id
           , first_name
           , last_name
           , performance_rating
       from employees
       where performance_rating ≥ 4.5) as employee_performance
join (select employee_id
           , count(project_id) as project_count
       from employee_projects
       group by employee_id) as project_allocations
using (employee_id)
where pa.project_count > 1
order by performance_rating desc
       , project_count desc
;
```

Oh no ! Subselects everywhere

- Each clause that require a *set of record*
 - can be replaced by a subselect
- Each clause that require a *value*
 - can be replaced by a subselect that yields only ONE element

Common Table Expression

```
with employee_performance as
    (select employee_id, first_name, last_name, performance_rating
     from employees
     where performance_rating ≥ 4.5 )

, project_allocations as
    (select employee_id, count(project_id) as project_count
     from employee_projects
     group by employee_id )

, high_performers_multiple_projects as
    (select first_name, last_name, performance_rating, project_count
     from employee_performance ep
     join project_allocations pa using (employee_id)
     where project_count > 1)

select *
from high_performers_multiple_projects
order by performance_rating desc, project_count desc;
```


Common Table Expression

- Extract logical subqueries from the main one
- Can even be recursive (but out of the scope of this talk)
- Warning : They can introduce memoization
 - `materialized/not materialized prefixes`

Views

```
create view employee_performance as
    select employee_id, first_name, last_name, performance_rating
    from employees
    where performance_rating ≥ 4.5;

create view project_allocations as
    select employee_id, count(project_id) as project_count
    from employee_projects
    group by employee_id;

select first_name
       , last_name
       , performance_rating
       , project_count
from employee_performance
     join project_allocations using (employee_id)
where project_counts > 1
order by (performance_rating, project_count) desc
;
```

Views

- Extract logical subqueries into reusable piece of code
- Allow neat tricks for schema upgrade !
- Warning : Documentation

Remember ?

```
select first_name
       , last_name
       , performance_rating
       , project_count
from (select employee_id
           , first_name
           , last_name
           , performance_rating
       from employees
       where performance_rating ≥ 4.5) as employee_performance
join (select employee_id
           , count(project_id) as project_count
       from employee_projects
       group by employee_id) as project_allocations
using (employee_id)
where pa.project_count > 1
order by performance_rating desc
       , project_count desc
;
```

Remember ?

```
select first_name
       , last_name
       , performance_rating
       , (select count(project_id)
          from employee_projects ep
         where employee_id = ep.employee_id) as project_count

from employees

where performance_rating ≥ 4.5
      and project_count > 1

order by performance_rating desc
       , project_count desc
;
```

Functions

```
create function get_project_count( employee_id int )  
returns bigint as $$  
begin return  
    select count(project_id)  
    from employee_projects ep  
    where employee_id = ep.employee_id;  
end;  
$$ language plpgsql;
```

Functions

```
select first_name
       , last_name
       , performance_rating
       , get_project_count(employee_id) as project_count

from employees

where performance_rating ≥ 4.5
      and project_count > 1

order by ( performance_rating, project_count ) desc
;
```

Functions

- In a functional language, using functions is a good idea
- As testable as any SQL query
- Composes way better than fragments

Set returning functions

```
create function project_stats( employee_id int )
returns table ( rd_project_count      bigint
               , it_project_count     bigint
               , pi_project_count     bigint
               , mean_project_duration interval) as $$
begin return query
    select count(project_id) filter (where type = 'R&D') as rd_project_count
          , count(project_id) filter (where type = 'IT')  as it_project_count
          , count(project_id) filter (where type = 'PI')  as pi_project_count
          , avg(age(start_at, end_at)) as mean_project_duration
    from employee_projects ep
    where employee_id = ep.employee_id;
end;
$$ language plpgsql;
```

SRF + Lateral

```
select first_name
       , last_name
       , performance_rating
       , rd_project_count
       , it_project_count
       , pi_project_count
       , mean_project_duration

from employees
     join lateral project_allocations(employee_id) on true

where performance_rating ≥ 4.5
       and project_count > 1

order by ( performance_rating, project_count ) desc
;
```

SRF + Lateral

```
select first_name
       , last_name
       , performance_rating
       , rd_project_count
       , it_project_count
       , pi_project_count
       , mean_project_duration

from employees
     , lateral project_allocations(employee_id)

where performance_rating ≥ 4.5
       and project_count > 1

order by ( performance_rating, project_count ) desc
;
```

Setup your development loop

Tools for composition

Step up your indexing game

Data types shenanigans

Indices 101

Persistent data structure to speed up data access:

- `btree` default index type.
 - Allows range lookup and sorting
- `brin` block range index.
 - Lighter but data must be in insert order
- `gin` inverted index.
 - Allows multi-key lookup but no sorting.
- `gist` framework to create complex datatype indices.
 - Allows distance based lookup and sorting

Indices 101

Stored in pages like tables : hence tablespace

```
create index ...  
tablespace nvme_drive;
```




Only one index can be used to scan in order

```
create table action (  
    id          bigserial    primary key,  
    user_id     uuid         not null,  
    ts          timestamptz  not null,  
    asset       text         not null,  
    status      text         not null  
);
```


Multicolumns indices

```
create index on action (user_id, ts desc)
```





Order of the columns is important !

-  user_id = ???
-  user_id = ??? and ts <= ???
-  ts between ??? and ???

Multicolumns indices

```
create index on action (user_id, ts desc)
```

Order of the columns is important !

-  order by user_id
-  order by user_id, ts desc
-  order by ts
-  order by ts desc, user_id

Included fields in index

```
create index on action (user_id, ts desc)
include (asset)
```

- Can make a query skip heap storage
- IndexOnlyScan vs IndexScan
- Useful when joining partial data

Partial Index

```
create index on action (user_id, ts desc)
include (asset)
where status = 'canceled'
```

- Reduce index size (and potentially increase performance)
- Can be combined with `unique` to enforce uniqueness on a subsets of rows

Clustered Tables

```
cluster table_name using idx_action_user_id_ts;
```

- Physically **reorders table rows** based on an **index**.
- Faster index-based scans (esp. range queries).

Clustered Tables

```
cluster table_name using idx_action_user_id_ts;
```

- **Table locked** during clustering (exclusive access).
- Needs **manual re-cluster**
- Can increase **write overhead** if frequently updated.

Setup your development loop

Tools for composition

Step up your indexing game

Data types shenanigans

Using domains

```
create domain invoice_total_amount as numeric(20,2)
check (
    value > 0 and value <= 1000000
)
constraint invoice_total_amount_range_check;
```


Using domains

```
create table invoice (  
    id          bigserial          primary key,  
    customer_id bigint            not null,  
    issued_at   date              not null,  
    total       invoice_total_amount not null  
);
```

Using complex data types : ranges & multiranges

- ranges are continuous intervals (e.g. dates, numbers)
- multiranges are sets of non-overlapping ranges
- Lots of operator / functions that handle all the tricky use cases
 - overlap
 - open / closed bounds
- Indexing support using gist and btree_gist

Using complex data types : ranges & multiranges

```
create table subscription (  
    id                bigserial primary key,  
    user_id           bigint      not null,  
    active_period     daterange  not null  
);
```

Using complex data types : ranges & multiranges

```
create index on subscription using gist (active_period);
```

Using complex data types : ranges & multiranges

```
select * from subscription  
where active_period @> date '2025-06-01';
```

```
select * from subscription  
where active_period -|- daterange('2025-06-01', '2025-07-01');
```

Using complex data types : exclude constraints

```
create table room_booking (  
    id                bigserial primary key,  
    room_id           bigint      not null,  
    booked_period      tsrange    not null,  
  
    exclude using gist (  
        room_id        with =,  
        booked_period  with &&  
    )  
);
```

Using complex data types : arrays

- collection of elements
- overhead for small collection is huge (24 bytes)
- Indexable using GIN (warning: no sorting!)
 - @> (inclusion)
 - && (overlap)
- For key-value, use `hstore`

Using complex data types : arrays and aggregations

```
with exploded(post_id, tag) as
    (select post_id, unnest(tags) from posts)

select post_id
    , array_agg(tag order by tag)
    filter (where tag <> 'draft')
    as cleaned
from exploded
group by post_id;
```


Using complex data types : ltree

- ltree stores hierarchical labels (path in a tree)
- lquery : pattern-matching syntax for querying ltrees
- indexable (GiST) on ltree and ltree[]

Using complex data types : ltree

```
create extension if not exists ltree;
```

```
create table product_category (  
    id    serial primary key,  
    name  text  not null,  
    path  ltree not null  
);
```

```
create index idx_product_path_gist  
on product_category  
using gist (path);
```

Using complex data types : ltree

```
-- 1. in Electronics
-- 2. Exactly 2-3 levels are left free
-- 3. Last label must be LED, OLED, or Speakers
-- 4. Exclude any category whose path contains Refurbished anywhere

select id, name, path
from   product_category
where  path ~ 'electronics.*{2,3}.{led|oled|speakers} & !*.refurbished';
```

JSONB : If your data has a schema : don't.

Using complex data types : jsonb

- Joins are fast.
- TOAST (The Oversized-Attribute Storage Technique)
 - Compressing and storing large field values (like long text or bytea) out-of-line in a separate pages to preserve row byte length
 - Reads are slower

Using complex data types : jsonb

- **base GIN**
 - `data @> '{"status": "active"}'` -- contains
 - `data ? 'email'` -- exists
- **path GIN (data jsonb_path_ops)**
 - only for @>
- **btree on extracted values**
 - `create index idx_data_price on my_table ((data ->> 'price'))`

JSONB : Great for aggregations

Using complex data types : jsonb aggregations

```
create table transactions (  
  transaction_id  uuid          primary key default gen_random_uuid(),  
  account_id     uuid          not null references accounts(account_id),  
  transaction_date timestampz not null  
);  
  
create table transaction_assets (  
  detail_id      uuid          primary key default gen_random_uuid(),  
  transaction_id uuid          not null references transactions(transaction_id),  
  asset_name     text          not null,  
  amount         numeric       not null  
);
```


Using complex data types : jsonb

```
create function get_transaction_transfers(transaction_id_input uuid)
returns jsonb
language plpgsql as $$
begin
    return select jsonb_agg(to_jsonb(*) - 'transaction_id')
           from transaction_assets
           where transaction_id = transaction_id_input;
end;
$$;
```

Using complex data types : jsonb

```
select transaction_id
       , transaction_date
       , get_transaction_transfers(transaction_id) as transfers
from transactions
where account_id = '<account_id>';
```

Using complex data types : cube

Builtin extension to manipulate *hyper-rectangles* and *points*

- spatial indexing
- Effectively multivariate order-independant index
- range-based searching.

Using complex data types : the cube's trick

```
create table customer_profile (  
    customer_id serial primary key,  
    logins integer not null,  
    session_duration numeric(5,2) not null,  
    purchases integer not null  
);
```

Using complex data types : the cube's trick

```
create index customer_behavior_idx
on customer_profile
using gist (
    cube(array[
        logins::float8,
        session_duration::float8,
        purchases::float8
    ])
);
```

Using complex data types : the cube's trick

```
select *  
from customer_profile  
where  
    cube(array[ logins::float8, session_duration::float8,  
purchases::float8])  
    <@ cube(array[10,20,1],array[50,40,10]);
```

Using complex data types : the cube's trick

QUERY PLAN

```
Bitmap Heap Scan on customer_profile  (cost=1.26..3.42 rows=2 width=24)
  Recheck Cond: (cube(ARRAY[...]) <@ '(10, 20, 1), (50, 40, 10)')::cube)
-> Bitmap Index Scan on customer_behavior_idx  (cost=0.00..1.26 rows=2 width=0)
     Index Cond: (cube(ARRAY[...]) <@ '(10, 20, 1), (50, 40, 10)')::cube)
```

What is first class Postgres ?

- Database is at the center of the system
- The schema is the source of truth
- Treat SQL as real code
- Commit to Postgres features

This meeting could have been an email.

This meeting could have been an email.

This service could have been a table.

